



Create S3 Buckets with Terraform



Nchindo Boris

The screenshot shows a web browser window with the URL developer.hashicorp.com/terraform/install. The page is titled "Install Terraform". On the left, there's a sidebar with "Install Terraform" selected. The main content area is for "macOS". It shows "Package manager" instructions: `brew tap hashicorp/tap` and `brew install hashicorp/tap/terraform`. Below that, "Binary download" sections are provided for "AMD64" and "ARM64" architectures, each with a "Download" button. To the right, there's an "About Terraform" section with a brief description and links to "Featured docs" like "Introduction to Terraform" and "Configuration Language". At the bottom, there's an "HCP Terraform" section with a "Try HCP Terraform for free" link.



Introducing Today's Project!

In this project, I will demonstrate how to use Terraform to automate the creation and management of an S3 bucket on AWS. The goal is to learn how Infrastructure as Code (IaC) can simplify cloud resource provisioning. By the end of this project, I will have a fully functioning Terraform setup that can create an S3 bucket, upload files to it, and manage it all through code.

Tools and concepts

Services I used were AWS S3, IAM, and Terraform. Key concepts I learnt include infrastructure as code (IaC), which involved tasks like; - Installing and configuring Terraform. - Configuring AWS credentials in the terminal. - Creating and managing S3 buckets with Terraform. - Uploading files to S3 bucket with Terraform.

Project reflection

This project took me approximately 45 minutes to complete. The most challenging part was setting up and troubleshooting AWS credentials to make sure Terraform could authenticate properly with my account. It was most rewarding to see the S3 bucket and the uploaded image in my AWS console



Nchindo Boris
NextWork Student

nextwork.org

I did this project today to build experience with Terraform and AWS, and to better understand how infrastructure as code works in a real-world workflow. It was a practical and rewarding way to connect concepts to cloud infrastructure.

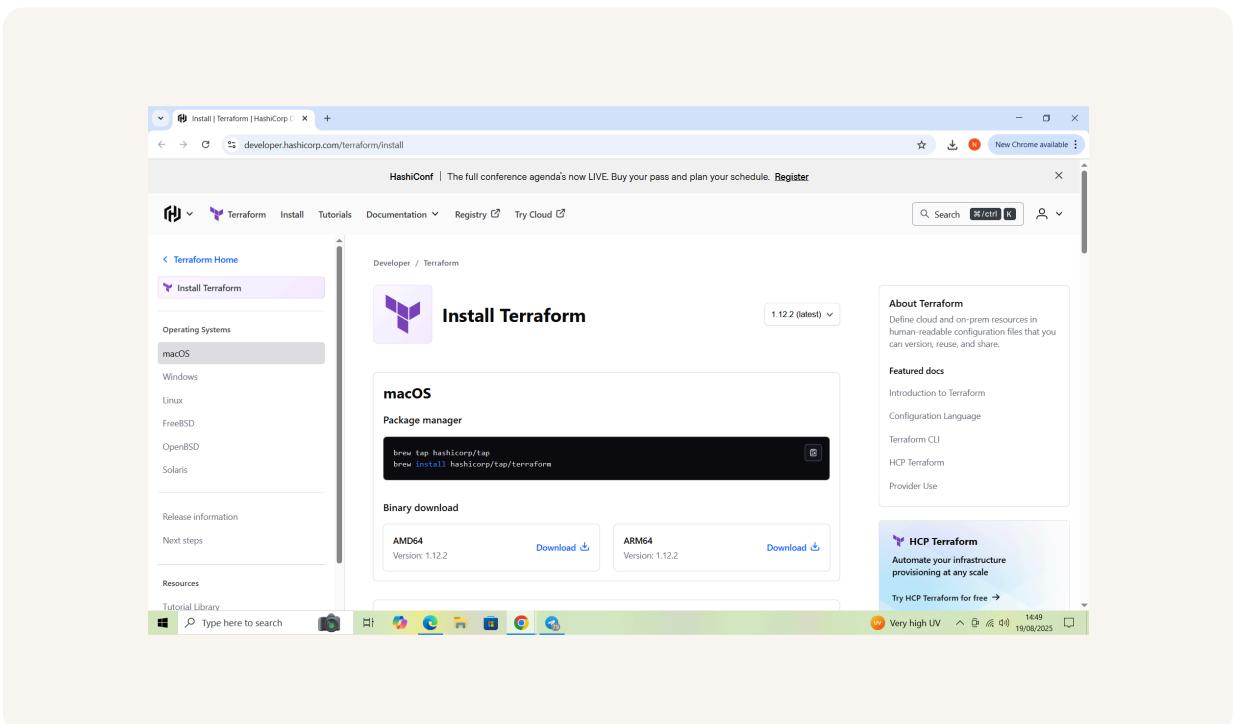


Introducing Terraform

Terraform is a tool that lets you build and manage infrastructure efficiently using code

Terraform is one of the most popular tools used for infrastructure as code (IaC), which is a practice that involves managing and provisioning cloud infrastructure using code instead of manually through a web console. Things built from code are built the same way every time. You can repair and rebuild things quickly, and other people can build identical instances of the same thing. This makes managing large-scale systems more efficient, less error-prone, and way faster than doing everything manually. Terraform is a top IaC tool because it supports multiple cloud providers, so you can set up multi-cloud infrastructures that use AWS, Azure, GCP and more. It also uses a simple configuration language.

Terraform uses configuration files to define and manage infrastructure. These files describe the desired state of your infrastructure, and Terraform figures out how to achieve that state. `main.tf` is a central file in a Terraform project.



Configuration files

The configuration is structured in blocks, such as provider and resource blocks. Each block has a specific role, like telling Terraform which cloud to use, what resources to create, or how to output useful information. The advantage of doing this is that it keeps your infrastructure organised, modular, and easy to maintain. You can update, reuse, or share parts of your setup without affecting unrelated components, making your infrastructure as code clean, scalable, and efficient.

My main.tf configuration has three blocks

The first block indicates that Terraform should use AWS as the cloud provider, specifying the region and credentials needed to interact with AWS services. The second block provisions an S3 bucket, telling AWS to create a new storage bucket and giving it a unique name. The third block manages the public access settings for the S3 bucket, ensuring that public access is fully blocked by setting all the appropriate security flags to true.



The screenshot shows a Microsoft code editor window titled "main.tf". The code is a Terraform configuration file:

```
C:/> Users > BORIS > Desktop > nextwork_terraform > main.tf
1 provider "aws" {
2   region = "us-east-2" # Update this to the Region closest to you
3 }
4
5 resource "aws_s3_bucket" "my_bucket" {
6   bucket = "nextwork-unique-bucket-sarika-9q" # Ensure this bucket name is globally unique
7 }
8
9 resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {
10   bucket = aws_s3_bucket.my_bucket.id
11   block_public_acls      = true
12   ignore_public_acls    = true
13   block_public_policy   = true
14   restrict_public_buckets = true
15 }
16 }
```

A tooltip at the bottom right of the editor window says: "You have Docker installed on your system. Do you want to install the recommended extensions from Microsoft for it?". The status bar at the bottom shows: Ln 6, Col 44, Spaces:4, UTF-8, CRLF, Plain Text, 1458, 28°C Mostly cloudy, 19/08/2025.

Customizing my S3 Bucket

For my project extension, I visited the official Terraform documentation to explore ways to customise the aws_s3_bucket resource further. The documentation shows detailed information on all the configurable properties for S3 buckets, such as bucket versioning, lifecycle rules, encryption options, logging, tags, and more. It provides examples and best practices on how to set these properties using Terraform code, enabling me to tailor the S3 bucket to meet specific needs beyond just creating a basic bucket.

I chose to customise my bucket by enabling versioning because it allows me to keep multiple versions of the objects stored in the bucket, which is useful for recovering from accidental deletes or overwrites. When I launch my bucket, I can verify my customisation by checking the bucket properties in the AWS S3 console or by running AWS CLI commands to see that versioning is enabled on the bucket. Additionally, I can upload and modify files to confirm that previous versions are being retained.

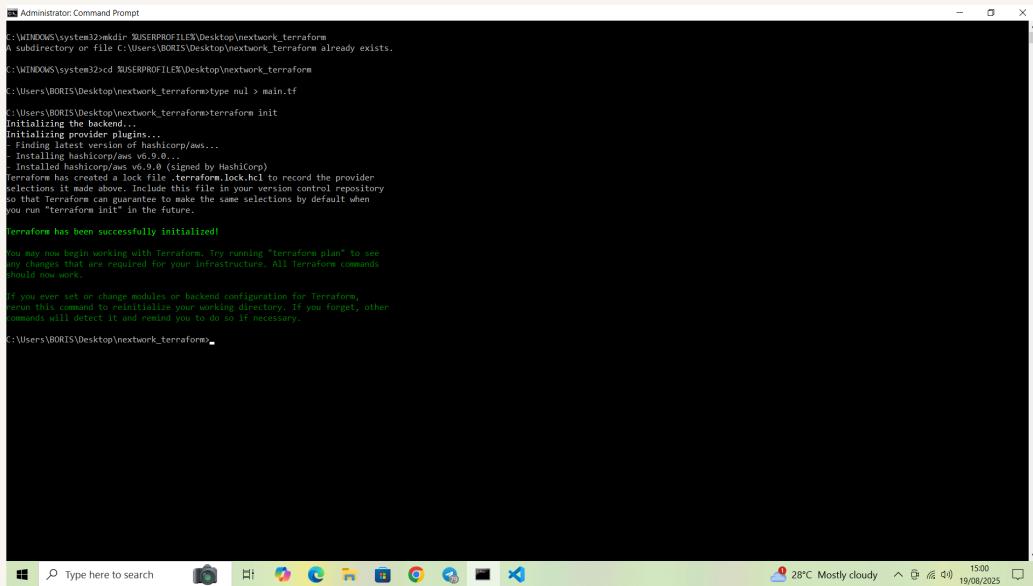


```
main.tf > ↵ resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block"
1 provider "aws" {
2   region = "your-region-code"
3 }
4
5 resource "aws_s3_bucket" "my_bucket" {
6   bucket = "nextwork-unique-bucket-johnnie-1699" # Make sure this bucket name is globally unique by typing a long random number here
7   tags = {
8     Name      = "My bucket"
9     Environment = "Test"
10    }
11  }
12
13 resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" [
14   bucket = aws_s3_bucket.my_bucket.id
15
16   block_public_acls      = true
17   ignore_public_acls     = true
18   block_public_policy    = true
19   restrict_public_buckets = true
20 ]
21
22 resource "aws_s3_bucket_versioning" "version_1" {
23   bucket = aws_s3_bucket.my_bucket.id
24   versioning_configuration {
25     status = "Enabled"
26   }
27 }
```

Terraform commands

I ran `terraform init` to set up my Terraform project by downloading the necessary provider plugins, configuring the backend to keep track of my infrastructure state, preparing any modules used in the configuration, and creating a lock file to ensure consistent versions across environments.

Next, I ran `terraform plan` to generate an execution plan that shows exactly what changes Terraform will make to my AWS infrastructure based on my configuration. This lets me review what resources will be created, updated, or destroyed before actually applying any changes.



```
C:\Windows\System32>mkdir %USERPROFILE%\Desktop\network_terraform
A subdirectory or file C:\Users\BORIS\Desktop\network_terraform already exists.

C:\Windows\System32>cd %USERPROFILE%\Desktop\network_terraform
C:\Users\BORIS\Desktop\network_terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/aws...
- Installing hashicorp/aws v0.9.0...
Installing provider dependencies... (hashicorp/aws v0.9.0)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
please take a moment to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\BORIS\Desktop\network_terraform>
```

AWS CLI and Access Keys

When I tried to plan my Terraform configuration, I received an error message that says my AWS credentials were missing or not configured because Terraform couldn't authenticate with AWS without proper access keys set up through the AWS CLI

To resolve my error, first, I installed AWS CLI, which is a command-line tool that allows me to manage AWS services directly from my terminal. This means I can run commands to configure and control AWS resources without using the web console, making it easier for Terraform to authenticate and interact with my AWS account.

I set up AWS access keys to allow the AWS CLI to securely authenticate and communicate with my AWS account. This way, Terraform can use the CLI with the proper permissions to create and manage resources like the S3 bucket without needing me to manually log in through the web console each time.



```
Planning failed. Terraform encountered an error while generating this plan.

Error: Retrieving AWS account details: validating provider credentials: retrieving caller identity from STS: operation error STS: GetCallerIdentity
      https response error StatusCode: 403, RequestID: 74f5c50d-708e-4e19-a3c4-28a28bc33bda, api error InvalidClientTokenId: The security token included in the request is invalid.

with provider["registry.terraform.io/hashicorp/aws"],
on main.tf line 1, in provider "aws":
  1: provider "aws" {
```

Lanching the S3 Bucket

I ran terraform apply to apply the changes written in the Terraform configuration. It creates, modifies, or deletes resources in the infrastructure based on code. Running terraform apply will affect my AWS account by creating the S3 bucket in my AWS account

The sequence of running terraform init, plan, and apply is crucial because each command prepares and validates the steps for deploying infrastructure safely and correctly: terraform init sets up the project by downloading necessary provider plugins (like AWS), initialising the backend, and preparing the environment to use Terraform. terraform plan then analyses the current configuration and compares it with the existing infrastructure (if any), showing you a preview of what changes Terraform will make. This helps avoid mistakes. terraform apply takes that plan and executes it, creating, updating, or deleting resources in your cloud environment. Running them in this order ensures that you initialise properly, preview your changes, and only then apply them, reducing the risk of accidental misconfigurations or resource loss.



The screenshot shows the AWS S3 console interface. The left sidebar has sections for General purpose buckets, Directory buckets, Table buckets, Vector buckets, Access Grants, Access Points (General Purpose Buckets, FSx file systems), Access Points (Directory Buckets), Object Lambda Access Points, Multi-Region Access Points, Batch Operations, and IAM Access Analyzer for S3. Below this is a section for Block Public Access settings for this account. Under Storage Lens, there are links for Dashboards, Storage Lens groups, and AWS Organizations settings.

The main content area displays the "General purpose buckets" tab, which lists one bucket named "nextwork-unique-bucket-suika-96". The bucket details show it was created in "US East (Ohio) us-east-2" on "August 19, 2025, 15:10:34 (UTC+01:00)". There are buttons for "Copy ARN", "Empty", "Delete", and "Create bucket". A search bar at the top of the list table allows searching by name. To the right of the bucket list are three cards: "Account snapshot" (updated daily, view dashboard), "Storage Lens provides visibility into storage usage and activity trends.", and "External access summary - new" (updated daily, info, external access findings help identify bucket permissions that allow public access or access from other AWS accounts).

The bottom of the screen shows the Windows taskbar with icons for CloudShell, Feedback, Start button, Task View, File Explorer, Edge, Google Chrome, File Manager, and Taskbar settings. The system tray shows the date (19/08/2025), time (15:11), weather (28°C Mostly cloudy), and battery status.



Uploading an S3 Object

I created a new resource block to upload an image file to my S3 bucket using Terraform. This block allows me to define the image as an object stored in the bucket, which means Terraform can manage not just the infrastructure but also the contents of that infrastructure. By including the file's name, location, and destination in the configuration, I ensure that every time I apply the Terraform plan, the object is reliably added or updated in the bucket.

We need to run `terraform apply` again because we made changes to our Terraform configuration; in this case, we added a new resource block to upload a file to our S3 bucket. Running `terraform apply` ensures those changes are applied to our AWS environment, keeping the actual infrastructure in sync with our updated configuration.

To validate that I've updated my configuration successfully, I checked my S3 bucket in the AWS Console and confirmed that the image file appeared there. Then, I downloaded the image from the bucket and verified that it matched the original file I uploaded, confirming that the upload worked.



```
C:\> Users > BORIS > Desktop > nextwork_terraform > main.tf
1 provider "aws" {
2   region = "us-east-2" # Update this to the Region closest to you
3 }
4
5 resource "aws_s3_bucket" "my_bucket" {
6   bucket = "nextwork-unique-bucket-svika-96" # Ensure this bucket name is globally unique
7 }
8
9 resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {
10   bucket = aws_s3_bucket.my_bucket.id
11
12   block_public_acls      = true
13   ignore_public_acls    = true
14   block_public_policy   = true
15   restrict_public_buckets = true
16 }
17
18 resource "aws_s3_object" "image" {
19   bucket = aws_s3_bucket.my_bucket.id # Reference the bucket ID
20   key    = "image.png" # Path in the bucket
21   source = "image.png" # Local file path
22 }
```



nextwork.org

The place to learn & showcase your skills

Check out nextwork.org for more projects

